


# Numerical Linear Algebra

- Useful in a wide range of applications, wherever physical systems can be described with linear equations – linear response systems, e.g.:
  - Reference frame transformations in both classical physics and special relativity.
  - Electric circuits, DC/AC (Kirchoff's laws).
  - Quantum mechanical states are vectors and linear combinations of vectors.
- Furthermore, linear algebra is often useful also in other numerical calculations. We will have various examples of this throughout this course (e.g. interpolation, least squares fitting, ODEs, ...)

# WMAP sky maps: linear algebra in action



**Map making requires several (very large) matrix multiplications and inversions.**

# Linear algebra

Computational linear algebra is a very wide field, we can cover only a part of it. Examples include:

solving systems of linear equations  $Ax=b$  for  $x$

solve the above for single  $A$  and many  $b$

compute inverse of matrix and determinant

different decompositions of matrices (e.g. LU, SVD)

dealing with over- and under-determined systems of equations, linear least-squares problem

computing eigenvalues and eigenvectors of matrices

“old” problem, large packages exist, e.g. LINPACK/LAPACK, much work has been done on optimisation and parallelization

# Linear algebra as benchmark

Q: How do we know which are the fastest, most powerful computers today?

Top 500 supercomputer

benchmark uses


LINPACK

implementation of

Gauss elimination  
w/partial pivoting.

(discussed in our  
next lecture)

→ Moore's law





# Linear systems of equations

$$a_{11}x_1 + a_{12}x_2 + \dots + a_{1n}x_n = b_1$$

$$a_{21}x_1 + a_{22}x_2 + \dots + a_{2n}x_n = b_2$$

$$\vdots \qquad \qquad \qquad \vdots$$

$$a_{n1}x_1 + a_{n2}x_2 + \dots + a_{nn}x_n = b_n$$

**or**

$$\begin{pmatrix} a_{11} & \cdots & a_{1n} \\ \vdots & \ddots & \vdots \\ a_{n1} & \cdots & a_{nn} \end{pmatrix} \cdot \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix} = \begin{pmatrix} b_1 \\ \vdots \\ b_n \end{pmatrix}$$

- Can be written in the matrix form  $Ax=b$

where:  $b$ =input,  $x$ =system response,  $A$ =physical system (independent of input).

- Solution exists if  $|A|$  is not 0 (number of independent equations = number of unknowns), i.e. when matrix  $A$  is non-singular.

# Linear systems of equations

$$a_{11}x_1 + a_{12}x_2 + \dots + a_{1n}x_n = b_1$$

$$a_{21}x_1 + a_{22}x_2 + \dots + a_{2n}x_n = b_2$$

$$\vdots \quad \quad \quad \vdots$$

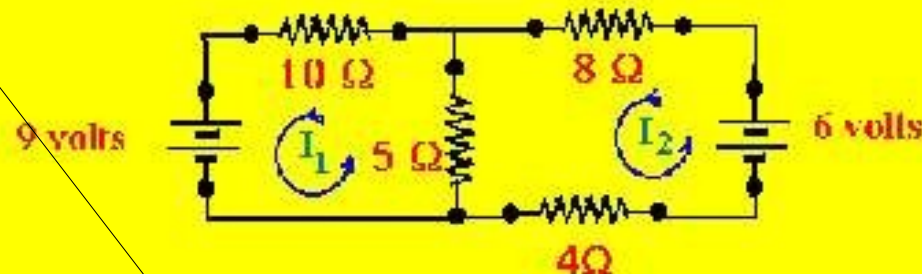
$$a_{n1}x_1 + a_{n2}x_2 + \dots + a_{nn}x_n = b_n$$

or

$$\begin{pmatrix} a_{11} & \cdots & a_{1n} \\ \vdots & \ddots & \vdots \\ a_{n1} & \cdots & a_{nn} \end{pmatrix} \cdot \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix} = \begin{pmatrix} b_1 \\ \vdots \\ b_n \end{pmatrix}$$

- Can be written in matrix form  $Ax = b$  where:  $b$ =input,  $x$ =output,  $A$ =system (independent equations)
- Solution exists if  $\text{rank}(A) = \text{rank}(b)$  and  $A$  is non-singular.

### Applying Kirchoff's Laws



1st Law:  $I_5 = I_1 - I_2$  = current through  $5\Omega$  resistor.

2nd Law:  $-4I_2 + 6 - 8I_2 + 5I_5 = 0$

2nd Law:  $-5I_5 - 10I_1 - 9 = 0$  **EXAMPLE**

Solve these three simultaneous equations to find the currents  $I_1$  -  $I_2$  and  $I_5$ .

# Solving linear systems

- Direct methods – transform the system into an equivalent one which is easier to solve by:
  - Exchanging two equations (changes sign of  $|A|$ ).
  - Multiplying an equation by a non-zero constant (multiplies  $|A|$  by the same constant).
  - Multiplying an equation by a nonzero constant and then subtracting it from another equation (leaves  $|A|$  unchanged).
- Iterative methods – start from a guess and then improve and refine it until convergence.

# Gauss elimination

We want to solve the linear system  $Ax=b$ , or:

$$\begin{array}{rcl} a_{11}x_1 + a_{12}x_2 + \dots + a_{1n}x_n & = & b_1 \\ a_{21}x_1 + a_{22}x_2 + \dots + a_{2n}x_n & = & b_2 \\ \vdots & & \vdots \\ a_{n1}x_1 + a_{n2}x_2 + \dots + a_{nn}x_n & = & b_n \end{array}$$

Gauss elimination: use first equation to remove the  $x_1$  term from all other equations by subtracting  $\lambda = a_{j1}/a_{11}$  times the first equation from the j-th equation ( $j > 1$ ):

$$\begin{aligned} a_{11}x_1 + a_{12}x_2 + \dots + a_{1n}x_n &= b_1 \\ a'_{22}x_2 + \dots + a'_{2n}x_n &= b'_2 \\ &\vdots \\ a'_{n2}x_2 + \dots + a'_{nn}x_n &= b'_n \end{aligned}$$



# Gauss elimination

This process is called “Gauss elimination”. We can then repeat it with the  $2^{\text{nd}}$  to  $n^{\text{th}}$  sub-matrix:

$$\text{Eq. } (i) \leftarrow \text{Eq. } (i) - \lambda \times \text{Eq. } (j)$$

Finally resulting in:

$$\begin{aligned} a_{11}x_1 + a_{12}x_2 + a_{13}x_3 + \dots + a_{1n}x_n &= b_1 \\ a'_{22}x_2 + a'_{23}x_3 + \dots + a'_{2n}x_n &= b'_2 \\ a''_{33}x_3 + \dots + a''_{3n}x_n &= b''_3 \\ &\vdots \\ a'''_{nn}x_n &= b'''_n \end{aligned}$$

→ we end up with an (upper) **triangular matrix** -- (forward) **elimination** phase, followed by **back-substitution** phase: solve last equation for  $x_n$  and work your way back up.

# Example

$$A = \begin{pmatrix} 1 & 2 & 3 \\ 2 & 2 & 3 \\ 3 & 3 & 3 \end{pmatrix} \quad b = \begin{pmatrix} 6 \\ 7 \\ 9 \end{pmatrix}$$

on board

# Gaussian elimination: algorithm

$$\left[ \begin{array}{cccccccc|c} A_{11} & A_{12} & A_{13} & \cdots & A_{1k} & \cdots & A_{1j} & \cdots & A_{1n} & b_1 \\ 0 & A_{22} & A_{23} & \cdots & A_{2k} & \cdots & A_{2j} & \cdots & A_{2n} & b_2 \\ 0 & 0 & A_{33} & \cdots & A_{3k} & \cdots & A_{3j} & \cdots & A_{3n} & b_3 \\ \vdots & \vdots & \vdots & & \vdots & & \vdots & & \vdots & \vdots \\ 0 & 0 & 0 & \cdots & A_{kk} & \cdots & A_{kj} & \cdots & A_{kn} & b_k \\ \hline \vdots & \vdots & \vdots & & \vdots & & \vdots & & \vdots & \vdots \\ 0 & 0 & 0 & \cdots & A_{ik} & \cdots & A_{ij} & \cdots & A_{in} & b_i \\ \vdots & \vdots & \vdots & & \vdots & & \vdots & & \vdots & \vdots \\ 0 & 0 & 0 & \cdots & A_{nk} & \cdots & A_{nj} & \cdots & A_{nn} & b_n \end{array} \right]$$

$$A_{ij} \leftarrow A_{ij} - \lambda A_{kj}, \quad j = k, k+1, \dots, n$$

$$b_i \leftarrow b_i - \lambda b_k$$

**Elimination phase**

# Gaussian elimination: algorithm

$$x_n = b_n / A_{nn}$$

**Back substitution phase**

$$x_k = \left( b_k - \sum_{j=k+1}^n A_{kj} x_j \right) \frac{1}{A_{kk}}, \quad k = n-1, n-2, \dots, 1$$



# Gauss Elimination: code

```
## module gaussElimin
''' x = gaussElimin(a,b) .
    Solves [a]{b} = {x} by Gauss elimination.
'''

from numpy import dot

def gaussElimin(a,b):
    n = len(b)
    # Elimination Phase
    for k in range(0,n-1):
        for i in range(k+1,n):
            if a[i,k] != 0.0:
                lam = a[i,k]/a[k,k]
                a[i,k+1:n] = a[i,k+1:n] - lam*a[k,k+1:n]
                b[i] = b[i] - lam*b[k]
    # Back substitution
    for k in range(n-1,-1,-1):
        b[k] = (b[k] - dot(a[k,k+1:n],b[k+1:n]))/a[k,k]
    return b
```

# Performance of Gauss elimination

Rough estimate:

forward elimination: 3 loops  $\rightarrow n^3$

back-substitution: 2 loops  $\rightarrow n^2$


More detailed analysis:  $n^3/3$  multiplications and  $n^2/2$  summations in total.

$\rightarrow O(n^3) \rightarrow n < 1000$  easy on a PC

For a general matrix no better scaling is possible – can only hope to improve pre-factor (and numerical stability)

Very brute-force methods may not work, e.g. Cramer's rule scales as  $n!$   $\rightarrow$  already  $n=20$  is numerically impossible!

# Linear Algebra: an Application

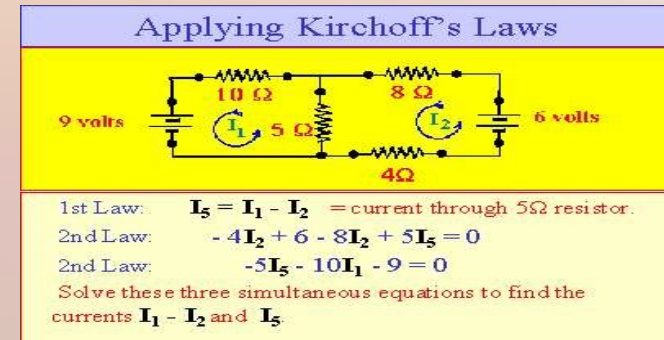


THE WORD "SUSTAINABLE" IS UNSUSTAINABLE.

# LU factorization

- Often we have to solve a linear system  $Ax=b$  for the same  $A$  (i.e. physical system), but different  $b$ 's (inputs).

Q: What does this correspond to in our Kirchhoff's Laws example?



- Solving the system again as a new one is possible, but is not the most efficient approach.
- Instead, we can re-use the expensive bit of the first solution, while getting the rest cheaply  $\rightarrow$  LU factorization:

$$\mathbf{L} = \begin{bmatrix} 1 & 0 & 0 \\ L_{21} & 1 & 0 \\ L_{31} & L_{32} & 1 \end{bmatrix} \quad \mathbf{U} = \begin{bmatrix} U_{11} & U_{12} & U_{13} \\ 0 & U_{22} & U_{23} \\ 0 & 0 & U_{33} \end{bmatrix}$$



# LU factorisation: solving the equations

Any square matrix  $A$  can be expressed as a product of a lower- and an upper-triangular matrix, such that  $A=LU$ . This is called **LU factorisation**. If  $A$  is nonsingular, so are  $L$  and  $U$ . If the split is available, obtaining the solution of  $Ax=b$  is done as follows:

- 1) split  $A=LU$  (the expensive bit,  $\sim n^3$  oper., independent of  $b$ )
- 2) solve  $Lz=b$  (cheap,  $\sim n^2$  operations)
- 3) solve  $Ux=z$ , (cheap,  $\sim n^2$  operations)

then,  $Ax=LUx=Lz=b$ , so we have the solution. The two triangular systems are easy to solve:

$l_{11}z_1=b_1 \rightarrow$  find  $z_1$ , then back-substitute in previous equations.  
Same for solving  $Ux=z$ , once we know  $z$ .

# LU factorisation: method

How do we do the LU factorization itself? We have already done it during Gauss elimination!

→ Just requires us to **store the Gauss multipliers  $\lambda$** .

U is just the usual upper triangular matrix from the forward Gauss elimination.

**L has 1 on diagonal and  $L_{ji} = \lambda_{ji}$**  where  $\lambda_{ji}$  is the multiplier required to eliminate  $A_{ji}$ . Just overwrite  $A_{ij}$  with  $\lambda$  instead of setting to zero in Gauss elimination algorithm.

$$\mathbf{L} = \begin{bmatrix} 1 & 0 & 0 \\ L_{21} & 1 & 0 \\ L_{31} & L_{32} & 1 \end{bmatrix} \quad \mathbf{U} = \begin{bmatrix} U_{11} & U_{12} & U_{13} \\ 0 & U_{22} & U_{23} \\ 0 & 0 & U_{33} \end{bmatrix} \quad \longrightarrow \quad [\mathbf{L} \setminus \mathbf{U}] = \begin{bmatrix} U_{11} & U_{12} & U_{13} \\ L_{21} & U_{22} & U_{23} \\ L_{31} & L_{32} & U_{33} \end{bmatrix}$$

# LU factorisation: expense

How many operations does LU factorization require?

Same as the Gaussian elimination!

Count them and remember that elimination is  $O(n^3)$ , while back-substitution is  $O(n^2)$ . We do same operation in each case, just in a different order.

Each additional solution of  $LUx=b$  for a new 'b' requires just  $O(n^2)$  operations, i.e. is relatively cheap (for large  $n$ , anyway).

# LU factorisation: notes and uses

The LU factorization of a matrix  $A$  is not unique, the one we discussed is called “Doolittle reduction”. Others are discussed in the book.

A special case: **for symmetric and positive-definite matrix**: can choose  $L = U'$ ,  $A = L L^T \rightarrow$  **Cholesky** factorisation.

Uses of LU:

**inverse**  $X$  of matrix  $A$ :  $\sum_j A_{ij} X_{jk} = \delta_{ik} \rightarrow$  think of  $X_{jk}$  as  $X_j(k)$  and  $\delta_{ik} = b_i(k)$   $k = 1, \dots, n$ , then solve  $LUx=b$  for  $n$  unit vector right-hand sides, then the solution vectors  $x$  are the columns of  $A^{-1}$  - matrix inverse is  $O(n^3)$  in operation count.


**determinant** of  $A$ : just product of diagonal elements  $U_{ii}$ :  
 $\det(A) = \det(L) \cdot \det(U) = \prod_i U_{ii}$



# Tri-diagonal matrices

- Some special (but important!) matrices can lead to systems of linear equations that are solved much faster than  $O(n^3)$ .
- An important example are **tridiagonal** matrices (which we will encounter later in this course in e.g. cubic spline interpolation and also in finite-difference DE solvers). These are a special case of so-called **sparse** matrices (i.e. matrices where a large number of values are zeros).
- LU decomposition preserves the banded structure:

$$A = \begin{bmatrix} X & X & 0 & 0 & 0 \\ X & X & X & 0 & 0 \\ 0 & X & X & X & 0 \\ 0 & 0 & X & X & X \\ 0 & 0 & 0 & X & X \end{bmatrix}$$



$$L = \begin{bmatrix} X & 0 & 0 & 0 & 0 \\ X & X & 0 & 0 & 0 \\ 0 & X & X & 0 & 0 \\ 0 & 0 & X & X & 0 \\ 0 & 0 & 0 & X & X \end{bmatrix}$$

$$U = \begin{bmatrix} X & X & 0 & 0 & 0 \\ 0 & X & X & 0 & 0 \\ 0 & 0 & X & X & 0 \\ 0 & 0 & 0 & X & X \\ 0 & 0 & 0 & 0 & X \end{bmatrix}$$

# Tri-diagonal matrices: storage

- Normally we store  $e$ ,  $d$ , and  $c$  as vectors, rather than the full matrix, which provides huge savings in storage.

$$A = \begin{bmatrix} d_1 & e_1 & 0 & 0 & \dots & 0 \\ c_1 & d_2 & e_2 & 0 & \dots & 0 \\ 0 & c_2 & d_3 & e_3 & \dots & 0 \\ 0 & 0 & c_3 & d_4 & \dots & 0 \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 0 & c_{n-1} & d_n \end{bmatrix}$$



$$c = \begin{bmatrix} c_1 \\ c_2 \\ \vdots \\ c_{n-1} \end{bmatrix} \quad d = \begin{bmatrix} d_1 \\ d_2 \\ \vdots \\ d_{n-1} \\ d_n \end{bmatrix} \quad e = \begin{bmatrix} e_1 \\ e_2 \\ \vdots \\ e_{n-1} \end{bmatrix}$$

- Example:

1000x1000 full matrix  $\rightarrow 10^6$  numbers

3-band only  $\rightarrow 2998$  numbers, **334x less!!**

# Tri-diagonal matrices: solution

- Standard methods like Gaussian elimination are easily adopted (but much faster) to this case. This can be solved by Gauss elimination with  $O(n)$  operations:

$$\text{row } k \leftarrow \text{row } k - (c_{k-1}/d_{k-1}) \times \text{row } (k-1), \quad k = 2, 3, \dots, n$$

$$d_k \leftarrow d_k - (c_{k-1}/d_{k-1})e_{k-1}$$

$$c_{k-1} \leftarrow c_{k-1}/d_{k-1}$$

**LU: Doolittle reduction**

**In Python:**

```
for k in range(1,n):  
    lam = c[k-1]/d[k-1]  
    d[k] = d[k] - lam*e[k-1]  
    c[k-1] = lam
```

$$A = \begin{bmatrix} d_1 & e_1 & 0 & 0 & \dots & 0 \\ c_1 & d_2 & e_2 & 0 & \dots & 0 \\ 0 & c_2 & d_3 & e_3 & \dots & 0 \\ 0 & 0 & c_3 & d_4 & \dots & 0 \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 0 & c_{n-1} & d_n \end{bmatrix}$$

**Then, the back substitution phase proceeds as before...**

# Tri-diagonal matrices: code from book

```
## module LUdecomp3
''' c,d,e = LUdecomp3(c,d,e): LU decomposition of tridiagonal matrix [c\d\e].
    On output{c},{d} and {e} are the diagonals of the decomposed matrix.
    x = LUsolve3(c,d,e,b): Solves [c\d\e]{x} = {b}, where {c}, {d} and {e}
    are the vectors returned from LUdecomp3.
'''
def LUdecomp3(c,d,e):
    n = len(d)
    for k in range(1,n):
        lam = c[k-1]/d[k-1]
        d[k] = d[k] - lam*e[k-1]
        c[k-1] = lam
    return c,d,e
def LUsolve3(c,d,e,b):
    n = len(d)
    for k in range(1,n):
        b[k] = b[k] - c[k-1]*b[k-1]
    b[n-1] = b[n-1]/d[n-1]
    for k in range(n-2,-1,-1):
        b[k] = (b[k] - e[k]*b[k+1])/d[k]
    return b
```



# Tridiagonal matrices: Scipy

- Naturally, there already exists an internal Scipy routine for solving banded problems:  
`scipy.linalg.solve_banded((l,u), cm, rhs)`
- The arguments are:
  - $(l,u) \rightarrow$  how many diagonals below and above the main diagonal – e.g. for tridiagonal matrix we have  $(l,u)=(1,1)$ .
  - `rhs`=a vector holding the right-hand side of the system
  - `cm` = a matrix storing (economically) the non-zero matrix elements, as follows:

$$\begin{pmatrix}
 1 & -5 & 3 & 0 & 0 & 0 & 0 & 0 \\
 3 & 2 & -3 & 5 & 0 & 0 & 0 & 0 \\
 0 & -2 & 1 & 5 & 9 & 0 & 0 & 0 \\
 0 & 0 & 9 & -1 & 5 & 4 & 0 & 0 \\
 0 & 0 & 0 & 0 & 2 & -3 & -2 & 0 \\
 0 & 0 & 0 & 0 & 2 & 0 & 1 & -6 \\
 0 & 0 & 0 & 0 & 0 & -3 & 2 & 7 \\
 0 & 0 & 0 & 0 & 0 & 0 & 9 & 1
 \end{pmatrix}
 \Rightarrow
 \begin{pmatrix}
 0 & 0 & 3 & 5 & 9 & 4 & -2 & -6 \\
 0 & -5 & -3 & 5 & 5 & -3 & 1 & 7 \\
 1 & 2 & 1 & -1 & 2 & 0 & 2 & 1 \\
 3 & -2 & 9 & 0 & 2 & -3 & 9 & 0
 \end{pmatrix}$$

# Gauss elimination with pivoting

Consider the following example:

$$\begin{pmatrix} 7 & -7 & 1 \\ -4 & 4 & 1 \\ 7 & 7 & -4 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} = \begin{pmatrix} 1 \\ 1 \\ 10 \end{pmatrix}$$

The algorithm already fails after first step, since  $\lambda = 14/0 \rightarrow$  division by zero! Similarly, a mix of very small and very large numbers will cause serious precision problems. Remedy: exchange 2nd and 3rd row  $\rightarrow$  “pivoting”. The diagonal element used to generate the division coefficient is called “pivot”.

Basic algorithm is numerically unstable, since a small pivot can lead to round-off and overflow. Instead, perform row interchanges (“partial pivoting”) so that you always place the largest element (on or below diagonal) on diagonal to be used as next pivot.

# Gauss with partial pivoting

Pivoting only requires a small change to the algorithm: in each row (after the `for i=1:n-1` statement), find  $p \geq i$  such that  $|A(p,i)| = \max(|A(j,i)|, j \geq i)$ , then interchange rows  $i$  and  $p$ .

There is also **full pivoting** which interchanges also columns, but partial pivoting is numerically nearly as stable, and is simpler, so we limit our discussion to it.

**Gauss elimination with partial pivoting** is simple to understand and implement, and reasonably stable and fast, and so **is a good test-method** if a more advanced “canned” routine gives suspect results.

# Gauss elimination with partial pivoting

```
## module gaussPivot
''' x = gaussPivot(a,b,tol=1.0e-9).
    Solves [a]{x} = {b} by Gauss elimination with scaled row pivoting
'''

from numpy import zeros, argmax, dot
import swap
import error

def gaussPivot(a,b,tol=1.0e-12):
    n = len(b)
    # Set up scale factors
    s = zeros(n)
    for i in range(n):
        s[i] = max(abs(a[i,:]))
    for k in range(0,n-1):
        # Row interchange, if needed
        p = argmax(abs(a[k:n,k])/s[k:n]) + k
        if abs(a[p,k]) < tol: error.err('Matrix is singular')
        if p != k:
            swap.swapRows(b,k,p)
            swap.swapRows(s,k,p)
            swap.swapRows(a,k,p)
```

# Gauss elimination with partial pivoting

The rest essentially the same as the simple Gaussian Elimination

```
# Elimination
```

```
    for i in range(k+1,n):
```

```
        if a[i,k] != 0.0:
```

```
            lam = a[i,k]/a[k,k]
```

```
            a[i,k+1:n] = a [i,k+1:n] - lam*a[k,k+1:n]
```

```
            b[i] = b[i] - lam*b[k]
```

```
if abs(a[n-1,n-1]) < tol: error.err('Matrix is singular')
```

```
# Back substitution
```

```
b[n-1] = b[n-1]/a[n-1,n-1]
```

```
for k in range(n-2,-1,-1):
```

```
    b[k] = (b[k] - dot(a[k,k+1:n],b[k+1:n]))/a[k,k]
```

```
return b
```

→ try this and the old version out on the previous example:

$$\begin{pmatrix} 7 & -7 & 1 \\ -4 & 4 & 1 \\ 7 & 7 & -4 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} = \begin{pmatrix} 1 \\ 1 \\ 10 \end{pmatrix}$$



# Ill-conditioned matrices

Some matrices are inherently close to being singular (i.e.  $\det(A)$  is close to 0) and lead to systems of equations that are difficult to solve.

Example: Hilbert matrix,  $h_{ij} = 1/(i+j-1)$

(III) conditioning can be quantified, but rules are complex. A simple 'rule of thumb' is that  $\det(A)$  should not be much smaller than the elements of  $A$ .

For ill-conditioned matrices even small changes of the coefficients can result in huge changes in the solution, which can make even small-ish systems of linear equations (e.g. 20) impossible to solve!

Fixes require special treatments, e.g. SVD (singular value decomposition) or extended precision routines.

# Comments

We discussed Gauss elimination and the need for (partial) pivoting.

This allows to solve  $Ax=b$  with forward elimination followed by back-substitution.

We also looked at LU decomposition, the inverse and the determinant of a matrix

These are direct methods, but for large sparse systems, iterative methods are often better (especially conjugate gradient methods and multigrid methods). We did not discuss the iterative methods here.

# Big libraries available

standard: **LAPACK** -> <http://www.netlib.org/lapack/>

- callable from Fortran, C, etc., implemented in SciPy Python package.

Generally, [netlib.org](http://www.netlib.org) is a good resource of information.

For C/C++, look at e.g. **GSL** (gnu scientific library) -> <http://www.gnu.org/software/gsl/>

Implementation “details” are very important for good performance.

→ commercial packages (e.g. NAG, IMSL) or

→ vendor implementations (intel MKL, AMD ACML)

# Python (NumPy/SciPy) commands

Internally, Python uses efficient **LAPACK** routines. Some are available in the NumPy module, and more are in the SciPy module ('import scipy', 'from scipy import linalg').

If A is a NumPy matrix and b is a vector, we have:

`x=solve(A,b)` (or `np.linalg.solve(A,b)`) – solves system  $Ax=b$  using Gauss elimination with partial pivoting (in the LU form) and row interchanges.

`linalg.lu(A)` – computes the LU factorisation explicitly.

`linalg.det(A)` – computes the determinant of A

`linalg.inv (A)` , or `A.I` – computes the inverse of A

`transpose(A)` – gives the transpose of A

Look them up in the help if you need them.